

[xcolor=dvipsnames]beamer
tikz,xcolor,soul media9 []beamerthemeMadrid
[shadow]beamerinnerthemeroounded []beamercolorthemerose
Theorem [theorem]Proposition [theorem]Corollary
enumerate graphicx amsmath,amssymb hyperref
[english]babel
[T1]fontenc ifthen latexsym,times,graphics,amssymb,fancybox,fancyhdr,bm
natbib
color
Sweave

Steps

1. *Import* data from various sources: the web, a database, a stored file.
2. *Clean* and format the data. Usually this means rows are observations and columns are variables.
3. *Analyze* the data using visualizations, modelling, or other methods.
4. *Communicate* your results.

¹Slide developed from G. Grolemund and H. Wickham.

Working with Data in R ¹

Steps

1. *Import* data from various sources: the web, a database, a stored file.
2. *Clean* and format the data. Usually this means rows are observations and columns are variables.
3. *Analyze* the data using visualizations, modelling, or other methods.
4. *Communicate* your results.

In this class, we learn tools and strategies for completing each step.

¹Slide developed from G. Grolemund and H. Wickham.

Function Definition

A function is a machine which turns input objects (*arguments*) into an output object (*return value*), according to a definite rule.

- Programming is writing functions to transform inputs into outputs easily and correctly.
- Good programming takes big transformations and breaks them down into smaller and smaller ones until you come to tasks which the built-in functions can do.

²Slide developed from C.R. Shalizi and A.C. Thomas (2014).

What is R?

What is R?

R is an open-source statistical programming software used by industry professionals and academics alike.

This means that R is supported by a community of users.

Will use R extensively in this class

- Download R at: <https://www.r-project.org>
- Download RStudio at: <https://www.rstudio.com>

What is R?

R is an open-source statistical programming software used by industry professionals and academics alike.

This means that R is supported by a community of users.

Will use R extensively in this class

- Download R at: <https://www.r-project.org>
- Download RStudio at: <https://www.rstudio.com>

Using R and RStudio

- The *editor* allows you to type and save code that you may want to reuse later.
- Basic interaction with R happens in the *console*. This is where you type R code.

Figure 1: Image of RStudio from G. Grolemund and H. Wickham

A Quick Example...

Type the following into your console:

```
> # Create a vector in R names "x"  
> x <- c(5, 29, 13, 87)  
> x
```

```
[1]  5 29 13 87
```

Two important ideas:

1. Commenting
2. Assignment

A Quick Example...

Type the following into your console:

```
> # Create a vector in R  
> x <- c(5, 29, 13, 87)  
> x
```

```
[1] 5 29 13 87
```

Two important ideas:

1. Commenting

- Anything after the # isn't evaluated by R.
- Used to leave notes for humans reading your code.
- Very important in our class. Comment your code!

2. Assignment

A Quick Example...

Type the following into your console:

```
> # Create a vector in R  
> x <- c(5, 29, 13, 87)  
> x
```

```
[1] 5 29 13 87
```

Two important ideas:

1. Commenting

2. Assignment

- The `<-` symbol means assign `x` the value `c(5, 29, 13, 87)`.
- Could use `=` instead of `<-` but this is discouraged.
- All assignments take the same form: `object_name <- value`.
- `c()` means “concatenate”.
- Type `x` into the console to print its assignment.

A Quick Example...

Type the following into your console:

```
> # Create a vector in R names "x"  
> x <- c(5, 29, 13, 87)  
> x
```

```
[1] 5 29 13 87
```

Note

The [1] tells us that 5 is the first element of the vector.

```
> # Create a vector in R names "x"  
> x <- 1:50  
> x
```

```
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18  
[19] 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36  
[37] 37 38 39 40 41 42 43 44 45 46 47 48 49 50
```

Variable Types, Vectors, & Matrices

R has a variety of variable types (or *modes*).

Modes

1. Numeric (3.7, 15907, 80.333)
2. Complex ($1 + 2i$)
3. Character ("Columbia", "Statistics is fun!", "HELLO WORLD")
4. Logical (TRUE, FALSE, 1, 0)

In this class, we are primarily concerned with numeric, character, and logical.

Let's check this out in R

'Numeric' variable type

```
> x <- 2  
> mode(x)
```

```
[1] "numeric"
```

```
> typeof(x)
```

```
[1] "double"
```

```
> y <- as.integer(3)  
> typeof(y)
```

```
[1] "integer"
```

Let's check this out in R

'Complex' variable type

```
> z <- 1 - 2i  
> z
```

```
[1] 1-2i
```

```
> typeof(z)
```

```
[1] "complex"
```


Let's check this out in R

'Character' variable type

```
> name <- "Columbia University"  
> name
```

```
[1] "Columbia University"
```

```
> typeof(name)
```

```
[1] "character"
```

Let's check this out in R

'Logical' variable type

```
> a <- TRUE  
> b <- F  
> a
```

```
[1] TRUE
```

```
> b
```

```
[1] FALSE
```

```
> typeof(a)
```

```
[1] "logical"
```

There are many data types in R.

Data Types

- Vectors
- Scalars
- Matrices
- Arrays
- Lists
- Dataframes

There are many data types in R.

Data Types

Vectors

- All elements must be the same type (mode).
- More to come in this lecture!

- Scalars
- Matrices
- Arrays
- Lists
- Dataframes

There are many data types in R.

Data Types

- Vectors

Scalars

- Treated as one-element vectors in R.
- Matrices
- Arrays
- Lists
- Dataframes

There are many data types in R.

Data Types

- Vectors
- Scalars

Matrices

- An array (rows and columns) of values.
- All values must be the same type (mode).
- More to come this lecture!

- Arrays
- Lists
- Dataframes

There are many data types in R.

Data Types

- Vectors
- Scalars
- Matrices

Arrays

- Similar to matrices, but with more than two dimensions.
- Lists
- Dataframes

There are many data types in R.

Data Types

- Vectors
- Scalars
- Matrices
- Arrays

Lists

- Like a vector, but elements can be of different modes.
 - Won't study lists explicitly, but encounter them all the time.
-
- Dataframes

There are many data types in R.

Data Types

- Vectors
- Scalars
- Matrices
- Arrays
- Lists

Dataframes

- Like a matrix, but elements can be of different modes.
- More to come next week!

Check Your Understanding

What mode are the following variables?

1. `3*TRUE?`
2. `"147"?`

Check Your Understanding

What mode are the following variables?

1. `3*TRUE`?
2. `"147"`?

Solutions

```
> 3*TRUE # Logicals in arithmetic
```

```
[1] 3
```

```
> mode(3*TRUE)
```

```
[1] "numeric"
```

```
> mode("147")
```

```
[1] "character"
```

Recall: Vectors

- Variable types are called modes.
- All elements of a vector are the same mode.
- Scalars are just single-element vectors.

Recall: Matrices

- All elements of a matrix are the same mode.
- A matrix is treated like a vector in R with two additional attributes: number of rows and number of columns.

Building Vectors in R

- Use the *concatenate* function `c()` to define a vector.

Some Examples

Defining a numeric vector:

```
> x <- c(2, pi, 1/2, 3^2)
> x
```

```
[1] 2.000000 3.141593 0.500000 9.000000
```

A character vector:

```
> y <- c("NYC", "Boston", "Philadelphia")
> y
```

```
[1] "NYC"           "Boston"        "Philadelphia"
```

Building Vectors in R

- The syntax `a:b` produces a *sequence* of integers ranging from *a* to *b*.
- The *repetition* function `rep(val, num)` repeats the value *val* a total of *num* times.

Some Examples

A sequential list of integers:

```
> z <- 5:10  
> z
```

```
[1] 5 6 7 8 9 10
```

Using `rep()` to create a 1's vector:

```
> u <- rep(1, 18)  
> u
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Building Vectors in R

- Alternately, could allocate space and then fill in element-wise.

Some Examples

```
> v <- c()  
> v[1] <- TRUE  
> v[2] <- TRUE  
> v[3] <- FALSE  
> v
```

```
[1] TRUE TRUE FALSE
```

Building Vectors in R

- The *concatenate* function `c()` can be nested.

Some Examples

```
> vec1 <- rep(-27, 3)
> vec1
```

```
[1] -27 -27 -27
```

```
> vec2 <- c(vec1, c(-26, -25, -24))
> vec2
```

```
[1] -27 -27 -27 -26 -25 -24
```


Building Matrices in R

- Use the function `matrix(values, nrow, ncol)` to define your matrix.
- In R, matrices are stored in *column-major order* (determines where the number go as in the following example).

Some Examples

Building a matrix that fills in by column:

```
> mat <- matrix(1:9, nrow = 3, ncol = 3)
> mat
```

	[,1]	[,2]	[,3]
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

Building Matrices in R

- Use the function `matrix(values, nrow, ncol)` to define your matrix.
- In R, matrices are stored in *column-major order* (determines where the number go as in the following example).

Some Examples

Building a matrix that fills in by row:

```
> new_mat <- matrix(1:9, nrow = 3, ncol = 3, byrow = TRUE)
> new_mat
```

	[,1]	[,2]	[,3]
[1,]	1	2	3
[2,]	4	5	6
[3,]	7	8	9

Building Matrices in R

- Alternately, could allocate space and fill in element-wise.
- Tell R the size of the matrix beforehand.

Some Examples

Allocating the space for a matrix then filling it in:

```
> this_mat <- matrix(nrow = 2, ncol = 2)
> this_mat[1,1] <- sqrt(27)
> this_mat[1,2] <- round(sqrt(27), 3)
> this_mat[2,1] <- exp(1)
> this_mat[2,2] <- log(1)
> this_mat
```

```
      [,1] [,2]
[1,] 5.196152 5.196
[2,] 2.718282 0.000
```

Building Matrices in R

- The *row bind* function `rbind()` also works, though it can be costly computationally. Similarly for *column bind* function `cbind()`.

Some Examples

```
> vec1 <- rep(0, 4)
> vec2 <- c("We're", "making", "matrices", "!")
> final_mat <- rbind(vec1, vec2)
> final_mat
```

	[,1]	[,2]	[,3]	[,4]
vec1	"0"	"0"	"0"	"0"
vec2	"We're"	"making"	"matrices"	"!"

Recall, matrix entries must all be the same type.

Building Matrices in R

- Name columns (rows) of a matrix using `colnames()` (`rownames()`).

Some Examples

```
> this_mat # Defined previously
```

```
      [,1] [,2]  
[1,] 5.196152 5.196  
[2,] 2.718282 0.000
```

```
> colnames(this_mat) # Nothing there yet
```

```
NULL
```

Building Matrices in R

- Name columns (rows) of a matrix using `colnames()` (`rownames()`).

Some Examples

```
> colnames(this_mat) <- c("Column1", "Column2")  
> this_mat
```

	Column1	Column2
[1,]	5.196152	5.196
[2,]	2.718282	0.000

Mixing Variable Modes

- When variable modes are mixed in vectors or matrices, R picks the 'least common denominator'.
- Use the *structure* function `str()` to display the internal structure of an R object.

Example

```
> vec <- c(1.75, TRUE, "abc")  
> vec
```

```
[1] "1.75" "TRUE" "abc"
```

```
> str(vec)
```

```
chr [1:3] "1.75" "TRUE" "abc"
```

Help in R

- Use a single question mark ? to get help about a specific function using form `?function name`.
 - Provides a description, lists the arguments (to the function), gives an example, etc.
- Use the double question mark to get help with a topic using form `??topic`.

How to get help in R

```
> # What does the str() function do?  
>  
> # Function help:  
> ?str  
> # Fuzzy matching:  
> ??"structure"
```


Code example.

Subsetting Vectors

- Use square brackets `[]` to extract elements or subsets of elements.

Example

```
> y <- c(27, -34, 19, 7, 61)
> y[2]
```

```
| [1] -34
```

```
> y[3:5]
```

```
| [1] 19  7 61
```

```
> y[c(1, 4)]
```

```
| [1] 27  7
```

Subsetting Vectors

- Use the same strategy to reassign elements of a vector.

Example

```
> y <- c(27, -34, 19, 7, 61)
> y
```

```
[1] 27 -34 19 7 61
```

```
> y[c(1, 4)] <- 0
> y
```

```
[1] 0 -34 19 0 61
```

Subsetting Vectors

- Negative values can be used to exclude elements.

Example

```
> y <- c(27, -34, 19, 7, 61)
> y
```

```
[1] 27 -34 19 7 61
```

```
> y[-c(1, 4)]
```

```
[1] -34 19 61
```

```
> y <- y[-1]
> y
```

```
[1] -34 19 7 61
```

Subsetting Matrices

- `mat[i,j]` returns the (i,j) th element of `mat`.
- `mat[i,]` returns the i^{th} row of `mat`.
- `mat[,j]` returns the j^{th} column of `mat`.

```
> mat <- matrix(1:8, ncol = 4)
> mat
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	3	5	7
[2,]	2	4	6	8

```
> mat[, 2:3]
```

	[,1]	[,2]
[1,]	3	5
[2,]	4	6

Subsetting Matrices

- Can use column names or row names to subset as well.
- Negative values are used to exclude elements.

```
> this_mat
```

```
      Column1 Column2  
[1,] 5.196152   5.196  
[2,] 2.718282   0.000
```

```
> this_mat[, "Column2"]
```

```
[1] 5.196 0.000
```

```
> this_mat[, -1]
```

```
[1] 5.196 0.000
```

An Extended Example: Image Data

But First...Packages!

What are packages?

- Packages are collections of functions, data, or code that extend the capabilities of base R.
- Some packages come pre-loaded but others must be downloaded and installed using function `install.packages("package name")`.
- An installed R package must be loaded in each session it is to be used using function `library("package name")`.

```
> # Installing the "pixmap" package.  
> install.packages("pixmap")  
> library("pixmap")
```


Image Data Example ³

- Images are made up of pixels which are arranged in rows and columns (like a matrix).
- Image data are matrices where each element is a number representing the intensity or brightness of the corresponding pixel.
- We will work with a greyscale image with numbers ranging from 0 (black) to 1 (white).

³Example developed from N. Matloff, “The Art of R Programming: A Tour of Statistical Software Design”.

Image Data Example (cont.)

```
> library(pixmap)
> casablanca_pic <- read.pnm("casablanca.pgm")
> casablanca_pic
```

Pixmap image

```
Type          : pixmapGrey
Size           : 360x460
Resolution     : 1x1
Bounding box   : 0 0 460 360
```

```
> plot(casablanca_pic)
```

Image Data Example (cont.)



Image Data Example (cont.)

```
> dim(casablanca_pic@grey)
```

```
[1] 360 460
```

```
> casablanca_pic@grey[360, 100]
```

```
[1] 0.4431373
```

```
> casablanca_pic@grey[180, 10]
```

```
[1] 0.9882353
```

Image Data Example (cont.)

Let's erase Rick from the image.

```
> casablanca_pic@grey[15:70, 220:265] <- 1
```

```
> plot(casablanca_pic)
```

Image Data Example (cont.)



Image Data Example (cont.)

- Use R's `locator()` function to find the rows and columns corresponding to Rick's face.
- A call to the function allows the user to click on a point in a plot and then the function returns the coordinates of the click.

Check Your Understanding

Using matrix `z`, what is the output of the following?

```
> z
```

	First	Second	Third
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

1. `z[2:3, "Third"]`?
2. `c(z[, -(2:3)], "abc")`?
3. `rbind(z[1,], 1:3)`?

Check Your Understanding

```
> z
```

	First	Second	Third
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

Solutions

```
> z[2:3, "Third"]
```

```
[1] 8 9
```

```
> c(z[, -(2:3)], "abc")
```

```
[1] "1" "2" "3" "abc"
```

Check Your Understanding

```
> z
```

	First	Second	Third
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

Solutions

```
> rbind(z[1,], 1:3)
```

	First	Second	Third
[1,]	1	4	7
[2,]	1	2	3

More with Vectors & Matrices and Linear Algebra Review

Reminder: Vector Algebra

Define vectors:

$$A = (a_1, a_2, \dots, a_N), \quad \text{and} \quad B = (b_1, b_2, \dots, b_N).$$

Then for c a scalar,

- $A + B = (a_1 + b_1, a_2 + b_2, \dots, a_N + b_N).$
- $cA = (ca_1, ca_2, \dots, ca_N).$
- Dot product: $A \cdot B = a_1b_1 + a_2b_2 + \dots + a_Nb_N.$
- Norm: $\|A\|^2 = A \cdot A = a_1^2 + a_2^2 + \dots + a_N^2.$

Reminder: Matrix Algebra

Define matrices:

$$A = \begin{pmatrix} a_1 & a_3 \\ a_2 & a_4 \end{pmatrix}, \quad \text{and} \quad B = \begin{pmatrix} b_1 & b_3 \\ b_2 & b_4 \end{pmatrix}.$$

Then for c a scalar,

- $A + B = \begin{pmatrix} a_1 + b_1 & a_3 + b_3 \\ a_2 + b_2 & a_4 + b_4 \end{pmatrix}.$
- $cA = \begin{pmatrix} ca_1 & ca_3 \\ ca_2 & ca_4 \end{pmatrix}.$
- Matrix Multiplication: $AB = \begin{pmatrix} a_1b_1 + a_3b_2 & a_1b_3 + a_3b_4 \\ a_2b_1 + a_4b_2 & a_2b_3 + a_4b_4 \end{pmatrix}.$

What if the dimensions of A and B are different?

Reminder: Matrix Operations

Define matrices:

$$A = \begin{pmatrix} a_{1,1} & a_{1,2} & \cdots & a_{1,m} \\ a_{2,1} & a_{2,2} & \cdots & a_{2,m} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n,1} & a_{n,2} & \cdots & a_{n,m} \end{pmatrix}, \quad \text{and } B = \begin{pmatrix} b_{1,1} & b_{1,2} \\ b_{2,1} & b_{2,2} \end{pmatrix}.$$

- The *transpose* of A is a $m \times n$ matrix:

$$t(A) = \begin{pmatrix} a_{1,1} & a_{2,1} & \cdots & a_{n,1} \\ a_{1,2} & a_{2,2} & \cdots & a_{n,2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1,m} & a_{2,m} & \cdots & a_{n,m} \end{pmatrix}.$$

- The *trace* of the square matrix B is the sum of the diagonal elements:
 $tr(B) = b_{1,1} + b_{2,2}.$

Reminder: Matrix Operations

Define matrices:

$$A = \begin{pmatrix} a_{1,1} & a_{1,2} & \cdots & a_{1,m} \\ a_{2,1} & a_{2,2} & \cdots & a_{2,m} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n,1} & a_{n,2} & \cdots & a_{n,m} \end{pmatrix}, \quad \text{and } B = \begin{pmatrix} b_{1,1} & b_{1,2} \\ b_{2,1} & b_{2,2} \end{pmatrix}.$$

- The *determinant* of square matrix B is $\det(B) = b_{1,1}b_{2,2} - b_{1,2}b_{2,1}$.
How do you find the determinant for an $n \times n$ matrix?
- The *inverse* of square matrix B is denoted B^{-1} and
 $BB^{-1} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$ and

$$B^{-1} = \frac{1}{\det(B)} \begin{pmatrix} b_{2,2} & -b_{1,2} \\ -b_{2,1} & b_{1,1} \end{pmatrix}.$$

How do you find the inverse of an $n \times n$ matrix?

Functions on Numeric Vectors

Useful R Functions

R function	Description
<code>length(x)</code>	Length of a vector <code>x</code>
<code>sum(x)</code>	Sum of a vector <code>x</code>
<code>mean(x)</code>	Arithmetic mean of a vector <code>x</code>
<code>quantiles(x)</code>	Sample quantiles of a vector <code>x</code>
<code>max(x)</code>	Maximum of a vector <code>x</code>
<code>min(x)</code>	Minimum of a vector <code>x</code>
<code>sd(x)</code>	Sample standard deviation of a vector <code>x</code>
<code>var(x)</code>	Sample variance of a vector <code>x</code>
<code>summary(x)</code>	Summary statistics of vector <code>x</code>

Reminder...

To access the help documentation of a known R function, use syntax
`?function`.

Example: Functions on Numeric Vectors

Example

To investigate the dependence of energy expenditure (y) on body build, researches used underwater weighing techniques to determine the fat-free body mass (x) for each of seven men. They also measured the total 24-hour energy expenditure for each man during conditions of quiet sedentary activity. The results are shown in the table.

Subject	1	2	3	4	5	6	7
x	49.3	59.3	68.3	48.1	57.61	78.1	76.1
y	1,894	2,050	2,353	1,838	1,948	2,528	2,568

```
> # Define covariate and response variable
> x <- c(49.3,59.3,68.3,48.1,57.61,78.1,76.1)
> y <- c(1894,2050,2353,1838,1948,2528,2568)
```

Example: Functions on Numeric Vectors (continued)

Example

Subject	1	2	3	4	5	6	7
x	49.3	59.3	68.3	48.1	57.61	78.1	76.1
y	1,894	2,050	2,353	1,838	1,948	2,528	2,568

```
> n <- length(x) # Sample size
```

```
> n
```

```
[1] 7
```

```
> max(x)
```

```
[1] 78.1
```

```
> sd(x)
```

```
[1] 12.09438
```

Example: Functions on Numeric Vectors (continued)

Example

Subject	1	2	3	4	5	6	7
x	49.3	59.3	68.3	48.1	57.61	78.1	76.1
y	1,894	2,050	2,353	1,838	1,948	2,528	2,568

```
> summary(x) # Summary statistics
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
48.10	53.46	59.30	62.40	72.20	78.10

```
> summary(y)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1838	1921	2050	2168	2440	2568

Element-Wise Operations for Vectors

Vectors x and y must have the same length. Let a be a scalar.

Element-Wise Operators

Operator	Description
$a + x$	Element-wise scalar addition
$a * x$	Element-wise scalar multiplication
$x + y$	Element-wise addition
$x * y$	Element-wise multiplication
$x ^ a$	Element-wise power
$a ^ x$	Element-wise exponentiation
$x ^ y$	Element-wise exponentiation

Recycling

Recall that a scalar is just a vector of length 1. When a shorter vector is added to a longer one, the elements in the shorter vectored are repeated. This is *recycling*.

Some Examples

```
> u <- c(1,3,5)
> v <- c(1,3,5)
> v + 4 # Recycling
```

```
[1] 5 7 9
```

```
> v + c(1,3) # Recycling
```

```
[1] 2 6 6
```

```
> v + u
```

```
[1] 2 6 10
```

Some Examples

Note: Operators are functions in R.

```
> u <- c(1,3,5)
> v <- c(1,3,5)
> '+'(v,u)
```

```
[1]  2  6 10
```

```
> '*'(v,u)
```

```
[1]  1  9 25
```

Line of Best Fit Example

Recall the energy expenditure versus fat-free body mass example.

Example


To investigate the dependence of energy expenditure (y) on body build, researchers used underwater weighing techniques to determine the fat-free body mass (x) for each of seven men. They also measured the total 24-hour energy expenditure for each man during conditions of quiet sedentary activity. The results are shown in the table.

Subject	1	2	3	4	5	6	7
x	49.3	59.3	68.3	48.1	57.61	78.1	76.1
y	1,894	2,050	2,353	1,838	1,948	2,528	2,568

Line of Best Fit Example (cont.)

Let's find the line of best fit.

```
> plot(x,y, xlab = "Body Mass", ylab = "Energy Expenditure")
```



Line of Best Fit Example (cont.)

Recall:

For the line of best fit, $\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$ where

$$\hat{\beta}_1 = \frac{S_{xy}}{S_{xx}} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}, \quad \text{and} \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}.$$

Solution:

```
> # First, compute x and y deviations
>   dev_x <- x - mean(x)
>   dev_y <- y - mean(y)
> # Next, compute sum of squares of xy and xx
>   Sxy <- sum(dev_x * dev_y)
>   Sxx <- sum(dev_x * dev_x)
```

Line of Best Fit Example (cont.)

Recall:

For the line of best fit, $\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$ where

$$\hat{\beta}_1 = \frac{S_{xy}}{S_{xx}} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}, \quad \text{and} \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}.$$

Solution:

```
> # Compute the estimated slope  
> Sxy/Sxx
```

```
[1] 25.01184
```

```
> # Compute the estimated intercept  
> mean(y) - (Sxy/Sxx) * mean(x)
```

```
[1] 607.6539
```

Functions for Numeric Matrices

Useful R Functions

R Function	Description
<code>A %*% B</code>	Matrix multiplication for compatible matrices A, B .
<code>dim(A)</code>	Dimension of matrix A .
<code>t(A)</code>	Transpose of matrix A .
<code>diag(x)</code>	Returns a diagonal matrix with elements x along the diagonal.
<code>diag(A)</code>	Returns a vector of the diagonal elements of A .
<code>solve(A,b)</code>	Returns x in the equation $b = Ax$.
<code>solve(A)</code>	Inverse of A where A is a square matrix.
<code>cbind(A,B)</code>	Combine matrices horizontally for compatible matrices A, B .
<code>rbind(A,B)</code>	Combine matrices vertically for compatible matrices A, B .

System of Linear Equations Example

Solve the system of equations:
$$\begin{cases} 3x - 2y + z = -1 \\ x + \frac{1}{2}y - 12z = 2 \\ x + y + z = 3 \end{cases}$$

Recall,

We can represent the system using matrices as follows:

$$\begin{pmatrix} 3 & -2 & 1 \\ 1 & \frac{1}{2} & -12 \\ 1 & 1 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix} = \begin{pmatrix} -1 \\ 2 \\ 3 \end{pmatrix}.$$

Then we would like to solve for vector (x, y, z) .

System of Linear Equations Example (cont.)

Recall,

$$\begin{pmatrix} 3 & -2 & 1 \\ 1 & \frac{1}{2} & -12 \\ 1 & 1 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix} = \begin{pmatrix} -1 \\ 2 \\ 3 \end{pmatrix}$$

Solution:

```
> # Define matrix A
> A <- matrix(c(3,1,1,-2,1/2,1,1,-12,3), nrow = 3)
> # Define vector b
> b <- c(-1, 2, 3)
> # Use the solve function
> solve(A, b)
```

```
[1] 1 2 0
```

System of Linear Equations Example (cont.)

Recall,

$$\begin{pmatrix} 3 & -2 & 1 \\ 1 & \frac{1}{2} & -12 \\ 1 & 1 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix} = \begin{pmatrix} -1 \\ 2 \\ 3 \end{pmatrix}$$

Let's use matrix multiplication to check that $\mathbf{x} = (1 \ 2 \ 0)^T$ is the correct solution to our system of equations.

Solution

```
> x <- c(1, 2, 0) # Define solution vector x
> A %*% x          # Then check with matrix multiplication
```

```
      [,1]
[1,]    -1
[2,]     2
[3,]     3
```

Element-wise Operations for Matrices

Let A and B be matrices of the same dimensions. Let a be a scalar.

Element-wise Operators

Operator	Description
$a + A$	Element-wise scalar addition
$a * A$	Element-wise scalar multiplication
$A + B$	Element-wise addition
$A * B$	Element-wise multiplication
$A ^ a$	Element-wise power
$a ^ A$	Element-wise exponentiation
$A ^ B$	Element-wise exponentiation

Eigenvalue Example

Check if 5 is an eigenvalue of the matrix $A = \begin{pmatrix} 1 & -2 \\ -2 & 4 \end{pmatrix}$.

Recall,

If λ is an eigenvalue of a square matrix A , then $Av = \lambda v$ for some non-zero vector v . Equivalently, if λ is a eigenvalue of A then $\det(A - \lambda I) = 0$ where I is an identity matrix.

Eigenvalue Example

Check if 5 is an eigenvalue of the matrix $A = \begin{pmatrix} 1 & -2 \\ -2 & 4 \end{pmatrix}$.

Solution:

```
> # Define matrix A
> A <- matrix(c(1, -2, -2, 4), nrow = 2, byrow = TRUE)
> # Define a 2 by 2 identity matrix
> identity <- diag(2)
> identity
```

	[,1]	[,2]
[1,]	1	0
[2,]	0	1

Eigenvalue Example

Check if 5 is an eigenvalue of the matrix $A = \begin{pmatrix} 1 & -2 \\ -2 & 4 \end{pmatrix}$.

Solution:

```
> # Check if 5 is an eigenvalue of A  
> det(A - 5*identity)
```

```
[1] 0
```

Filtering

Logical and Relational Operators

Logical Operator	Description
!	Negation (NOT)
&	AND
	OR

Relational Operator	Description
<, >	Less than, greater than
<=, >=	Less than or equal to, greater than or equal to
==	Equal to
!=	Not equal to

Examples of Logical and Relational Operators

Some Basic Examples

```
> 1 > 3
```

```
[1] FALSE
```

```
> 1 == 3
```

```
[1] FALSE
```

```
> 1 != 3
```

```
[1] TRUE
```

Examples of Logical and Relational Operators

Some Basic Examples

```
> (1 > 3) & (4*5 == 20)
```

```
[1] FALSE
```

```
> (1 > 3) | (4*5 == 20)
```

```
[1] TRUE
```

Examples of Logical and Relational Operators

Some Basic Examples

```
> c(0,1,4) < 3
```

```
[1] TRUE TRUE FALSE
```

```
> which(c(0,1,4) < 3)
```

```
[1] 1 2
```

```
> which(c(TRUE, TRUE, FALSE))
```

```
[1] 1 2
```

Examples of Logical and Relational Operators

Some Basic Examples

```
> c(0,1,4) >= c(1,1,3)
```

```
[1] FALSE TRUE TRUE
```

```
> c("Cat","Dog") == "Dog"
```

```
[1] FALSE TRUE
```


Filtering Examples

Sometimes we would like to extract elements from a vector or matrix that satisfy certain criteria.

Extracting Elements from a Vector

```
> w <- c(-3, 20, 9, 2)
> w[w > 3] ### Extract elements of w greater than 3
```

```
[1] 20  9
```

```
> ### What's going on here?
> w > 3
```

```
[1] FALSE  TRUE  TRUE FALSE
```

```
> w[c(FALSE, TRUE, TRUE, FALSE)]
```

```
[1] 20  9
```

Filtering Examples

```
> w <- c(-3, 20, 9, 2)
> ### Extract elements of w with squares between 3 and 10
> w[w*w >= 3 & w*w <= 10]
```

```
[1] -3  2
```

```
> w*w >= 3 ### What's going on here?
```

```
[1] TRUE TRUE TRUE TRUE
```

```
> w*w <= 10
```

```
[1] TRUE FALSE FALSE TRUE
```

```
> w*w >= 3 & w*w <= 10
```

```
[1] TRUE FALSE FALSE TRUE
```

Filtering Examples

Extracting Elements from a Vector

```
> w <- c(-1, 20, 9, 2)
> v <- c(0, 17, 10, 1)
> ### Extract elements of w greater than elements from v
> w[w > v]
```

```
[1] 20  2
```

```
> ### What's going on here?
> w > v
```

```
[1] FALSE  TRUE FALSE  TRUE
```

```
> w[c(FALSE, TRUE, FALSE, TRUE)]
```

```
[1] 20  2
```

Filtering Examples

Filtering Elements of a Matrix

```
> M <- matrix(c(rep(4,5), 5:8), ncol=3, nrow=3)
> M
```

	[,1]	[,2]	[,3]
[1,]	4	4	6
[2,]	4	4	7
[3,]	4	5	8

```
> ### We can do element-wise comparisons with matrices too.
> M > 5
```

	[,1]	[,2]	[,3]
[1,]	FALSE	FALSE	TRUE
[2,]	FALSE	FALSE	TRUE
[3,]	FALSE	FALSE	TRUE

Filtering Examples

```
> M
```

	[,1]	[,2]	[,3]
[1,]	4	4	6
[2,]	4	4	7
[3,]	4	5	8

```
> M[,3] < 8
```

```
[1] TRUE TRUE FALSE
```

```
> M[M[,3] < 8, ]
```

	[,1]	[,2]	[,3]
[1,]	4	4	6
[2,]	4	4	7

Filtering Examples

Reassigning Elements of a Matrix

```
> M
```

	[,1]	[,2]	[,3]
[1,]	4	4	6
[2,]	4	4	7
[3,]	4	5	8

```
> ### Assign elements greater than 5 with zero  
> M[M > 5] <- 0  
> M
```

	[,1]	[,2]	[,3]
[1,]	4	4	0
[2,]	4	4	0
[3,]	4	5	0

Check Your Understanding

Using matrix `z`, what is the output of the following?

```
> z
```

	First	Second	Third
[1,]	1	1	9
[2,]	2	0	16
[3,]	3	1	25

1. `z[z[, "Second"],]?`
2. `z[, 1] != 1?`
3. `z[(z[, 1] != 1), 3]?`

Check Your Understanding

```
> z
```

	First	Second	Third
[1,]	1	1	9
[2,]	2	0	16
[3,]	3	1	25

Solutions

```
> z[z[, "Second"], ]
```

	First	Second	Third
[1,]	1	1	9
[2,]	1	1	9

Check Your Understanding

```
> z
```

	First	Second	Third
[1,]	1	1	9
[2,]	2	0	16
[3,]	3	1	25

Solutions

```
> z[, 1] != 1
```

```
[1] FALSE TRUE TRUE
```

```
> z[(z[, 1] != 1), 3]
```

```
[1] 16 25
```

A Quick Note

```
> z
```

	First	Second	Third
[1,]	1	1	9
[2,]	2	0	16
[3,]	3	1	25

```
> z[(z[, 1] != 1), 3]
```

```
[1] 16 25
```

```
> z[(z[, 1] != 1), 3, drop = FALSE]
```

	Third
[1,]	16
[2,]	25

NA and NULL Values

NA and NULL

- NA indicates a missing value in a dataset.
- NULL is a value that doesn't exist and is often returned by expressions and functions whose value is undefined.

Example

```
> length(c(-1, 0, NA, 5))
```

```
[1] 4
```

```
> length(c(-1, 0, NULL, 5))
```

```
[1] 3
```

NA and NULL

Example

```
> ### Use na.rm = TRUE to remove NA values  
> t <- c(-1,0,NA,5)  
> mean(t)
```

```
[1] NA
```

```
> mean(t, na.rm = TRUE)
```

```
[1] 1.333333
```

```
> ### NA values are missing, but NULL values don't exist.  
> s <- c(-1, 0, NULL, 5)  
> mean(s)
```

```
[1] 1.333333
```

NA and NULL

NULL can be used is to build a vector in the following way:

```
> # Define an empty vector
>     x <- NULL
> # Fill in the vector
>     x[1] <- "Blue"
>     x[2] <- "Green"
>     x[3] <- "Red"
>     x
```

```
[1] "Blue" "Green" "Red"
```

- NULL is commonly used to build vectors in loops with each iteration adding another element.
- Filling in pre-allocated space is less expensive (computationally) than adding an element at each step.
- Loops will be introduced next lecture.

A Note on Lists

A list structure combines objects of different modes.

Recall, in vectors and matrices all elements must have the same mode.

To define a list:

- Use the function “list()”:

```
list(name1 = object1, name2 = object2, ...)
```

List component names (called [tags](#)) are optional.

Line of Best Fit Example

Recall, the energy expenditure versus fat-free body mass example one more time.

Example

To investigate the dependence of energy expenditure (y) on body build, researches used underwater weighing techniques to determine the fat-free body mass (x) for each of seven men. They also measured the total 24-hour energy expenditure for each man during conditions of quiet sedentary activity. The results are shown in the table.

Subject	1	2	3	4	5	6	7
x	49.3	59.3	68.3	48.1	57.61	78.1	76.1
y	1,894	2,050	2,353	1,838	1,948	2,528	2,568

Let's make a list of the values we've calculated for this example.

```
> # Combine data into single matrix
> data <- cbind(x, y)
> # Summary values for x and y
> sum_x <- summary(x)
> sum_y <- summary(y)
> # We computed Sxy and Sxx previously
> est_vals <- c(Sxy/Sxx, mean(y) - Sxy/Sxx*mean(x))
```

```
> # Define a list with different objects for each element
> body_fat <- list(variable_data = data,
+                  summary_x = sum_x, summary_y = sum_y,
+                  LOBF_est = est_vals)
```

Extracting Components of Lists

Extract an individual component “c” from a list called `lst` in the following ways:

- `lst$c`
- `lst[[i]]` where “c” is the i^{th} component.
- `lst[["c"]]`

Extracting Components of Lists

Energy expenditure versus fat-free body mass example

```
> # Extract the first list element  
> body_fat[[1]]
```

	x	y
[1,]	49.30	1894
[2,]	59.30	2050
[3,]	68.30	2353
[4,]	48.10	1838
[5,]	57.61	1948
[6,]	78.10	2528
[7,]	76.10	2568

Energy expenditure versus fat-free body mass example

```
> # Extract the Line of Best Fit estimates  
> body_fat$LOBF_est
```

```
[1] 25.01184 607.65386
```

```
> # Extract the summary of x  
> body_fat[["summary_x"]]
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
48.10	53.46	59.30	62.40	72.20	78.10